Software Development Practices for Public Health Bioinformatics

Week 04: Workflow Managers and Nextflow

A Mid-Atlantic Workforce Development Offering Provided by the Division of Consolidated Laboratory Services in Collaboration with Theiagen Genomics

Course Introduction

Training Workshop Instructors

Andrew Hale, BSc

- Bioinformatics Developer at Theiagen Genomics since 2025
- BSc in Computational Bioinformatics

Course Agenda

Course Agenda

Use of Workflow Managers

Week 4 - February 25/27, 2025

- Basics of workflow managers, with a focus on Nextflow
    - Package managers
    - Workflow managers
    - Nextflow

Package Managers

CONDA

Conda is a powerful command line tool for **package and environment management** that runs on Windows, macOS, and Linux.

Anaconda - Comes with a lot of pre-installed packages

Miniconda - Comes with the bare necessities

CONDA

MINICONDA®
= conda
+Python.exe
+base modules

ANACONDA®
= miniconda
+150 modules
+ user interface

CONDA

Why Conda?

- Non-Root access   ⟶   ● Use non-root user
- Dependency conflicts   ⟶   ● Installs in separate environments
- Shareable   ⟶   ● Installs are modular and easy to share
- Easy to execute   ⟶   ● Easy to load/unload environments

```
# Download Miniconda installer
wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh

# Run installer
bash Miniconda3-latest-Linux-x86_64.sh
```

Conda Rules of Thumb

There is an exception to this rule of thumb though!

Rule Number 1 - Keep base clean
Keep your base environment clean! Try to avoid installing anything in your base environment. If your base environment breaks, you have to reinstall Miniconda.

Rule Number 2 - Create environments
`conda create` is your friend, use it for everything! Treat environments as consumables, create them, install in them, delete them.
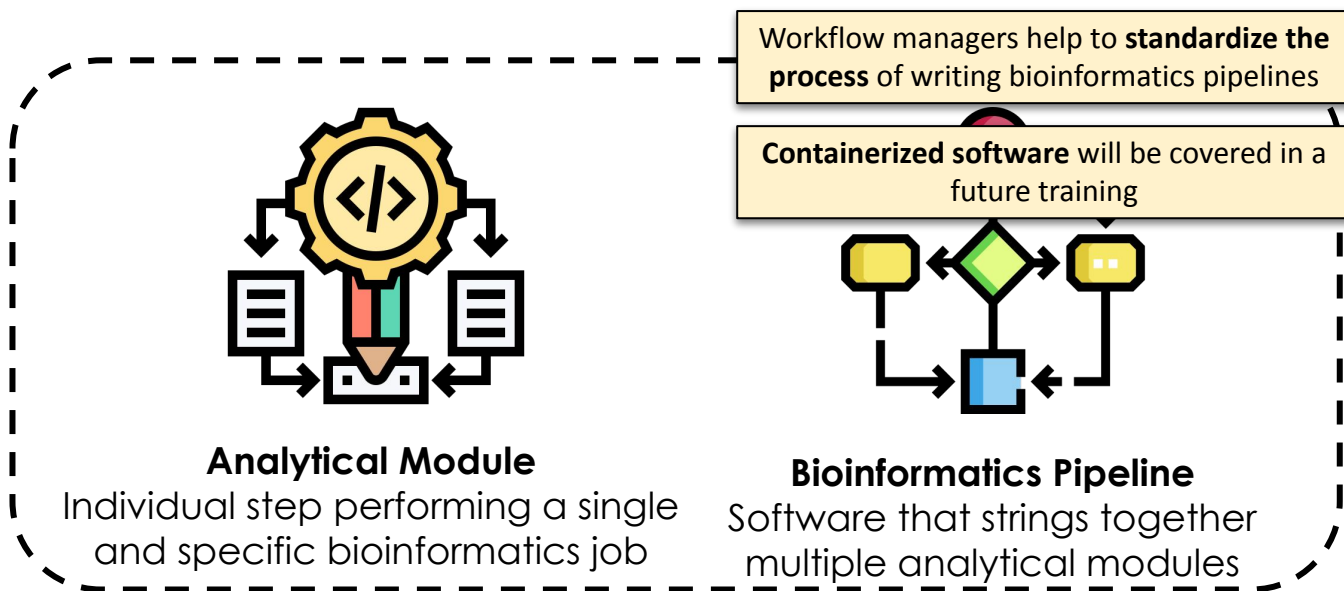
Rule Number 3 - Use containers for critical tasks
Conda is great, but for critical things its best to use containers (Docker or Singularity).

Workflow Managers

Introduction to Workflow Managers

Workflow managers provide a **standardized framework** for creating reproducible and interoperable **bioinformatics pipelines**, especially when **containerized software** are utilized

Workflow managers help to **standardize the process** of writing bioinformatics pipelines

Containerized software will be covered in a future training

Analytical Module
Individual step performing a single and specific bioinformatics job

Bioinformatics Pipeline
Software that strings together multiple analytical modules

Theiagen®

Introduction to Workflow Managers

Some workflow languages can be executed by multiple workflow engines

Workflow Language
Programming language used to describe the bioinformatics pipeline
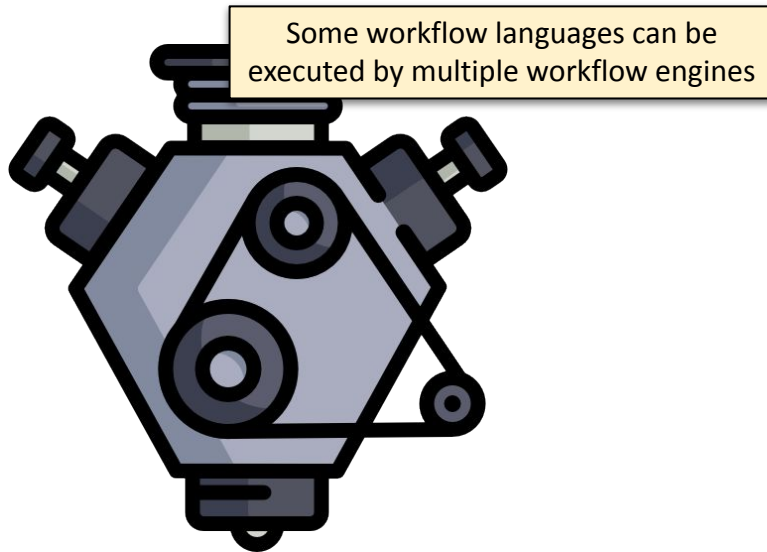
Workflow Engine
Software to interpret and execute the workflow itself

14

Nextflow

Introduction to Nextflow

Workflows
Using computers to collect, store, analyze and disseminate data and information

Large files
> 10 GB for a metagenomic sample
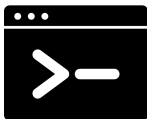
Many languages
Bash, Python, PERL…

Complex interactions
Network of software and their dependencies

Introduction to Nextflow

Nextflow
Managing modern workflows is complicated

Introduction to Nextflow

Reactive workflow framework

Create pipelines with asynchronous data streams

Programing DSL

Has its own language for building a pipeline

Containerized

Out of the box integrating with container engines

Simplifies the deployment of complex parallel and reactive workflows with resumability functionality

```
process sayHello {
  input:
    val cheers
  output:
    stdout

  """
  echo $cheers
  """
}

workflow {
  channel.of('Ciao','Hello','Hola') | sayHello | view
}
```

Theiagen

Introduction to Nextflow

Introduction to Nextflow
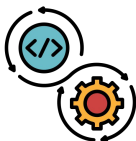
1. Nextflow Basics

Nextflow Basic Structure
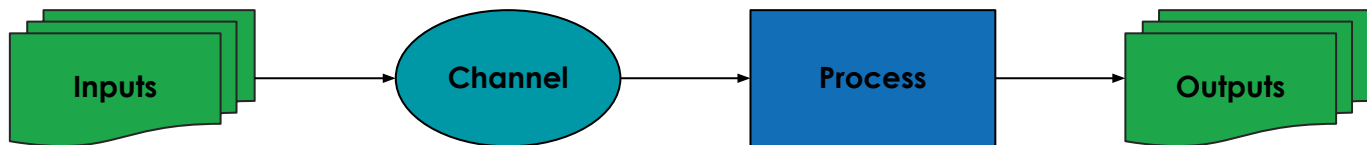
- Channel
  - Queue Channel
    - "First-in, First-out" (FIFO) connecting processes and operators
  - Value Channel
    - Stores a single value (e.g., genome size)
    - Can be a named list (e.g., ["id": "my_sample", "genome_size": 360000])

- Channels can be used by **Operators** and **Processes**

Nextflow Basic Structure

- Methods to **connect** channels, or **transform** values of channels
  - Filtering: filter, randomSample, take, unique
  - Transforming: collect, groupTuple, reduce
  - Text Processing: splitCsv, splitJson, splitText
  - Combining: combine, concat, join, mix
  - Forking: branch, multiMap
  - Math: count, max, min, sum
  - Other: ifEmpty, map, set, view
- More than 50 Operators available to use

Nextflow Basic Structure

Process example:

```
bwa mem reference.fa sample.fq \
  | samtools sort -o sample.bam
```

Nextflow Basic Structure

Process example:

```
process align_sample {

    input:
    file 'reference.fasta'
    file 'sample.fq'

    output:
    file 'sample.bam'

    script:
    """
    bwa mem reference.fa sample.fq \
        | samtools sort -o sample.bam
    """
}
```

Nextflow Basic Structure

Processes:

- The basic unit for executing user scripts
- [30+ directives](#) adjust optional settings
- Can be dynamic
- Inputs and outputs are channels
- Can be optional
- Conditional executions using 'when'
- The script block executes user code

```
process < name > {

  [ directives ]

  input:
   < process inputs >

  output:
   < process outputs >

  when:
   < condition >

  [script|shell|exec]:
  < user script to be executed >

}
```

Nextflow Basic Structure

FIFO Can Cause "Unexpected" Results:
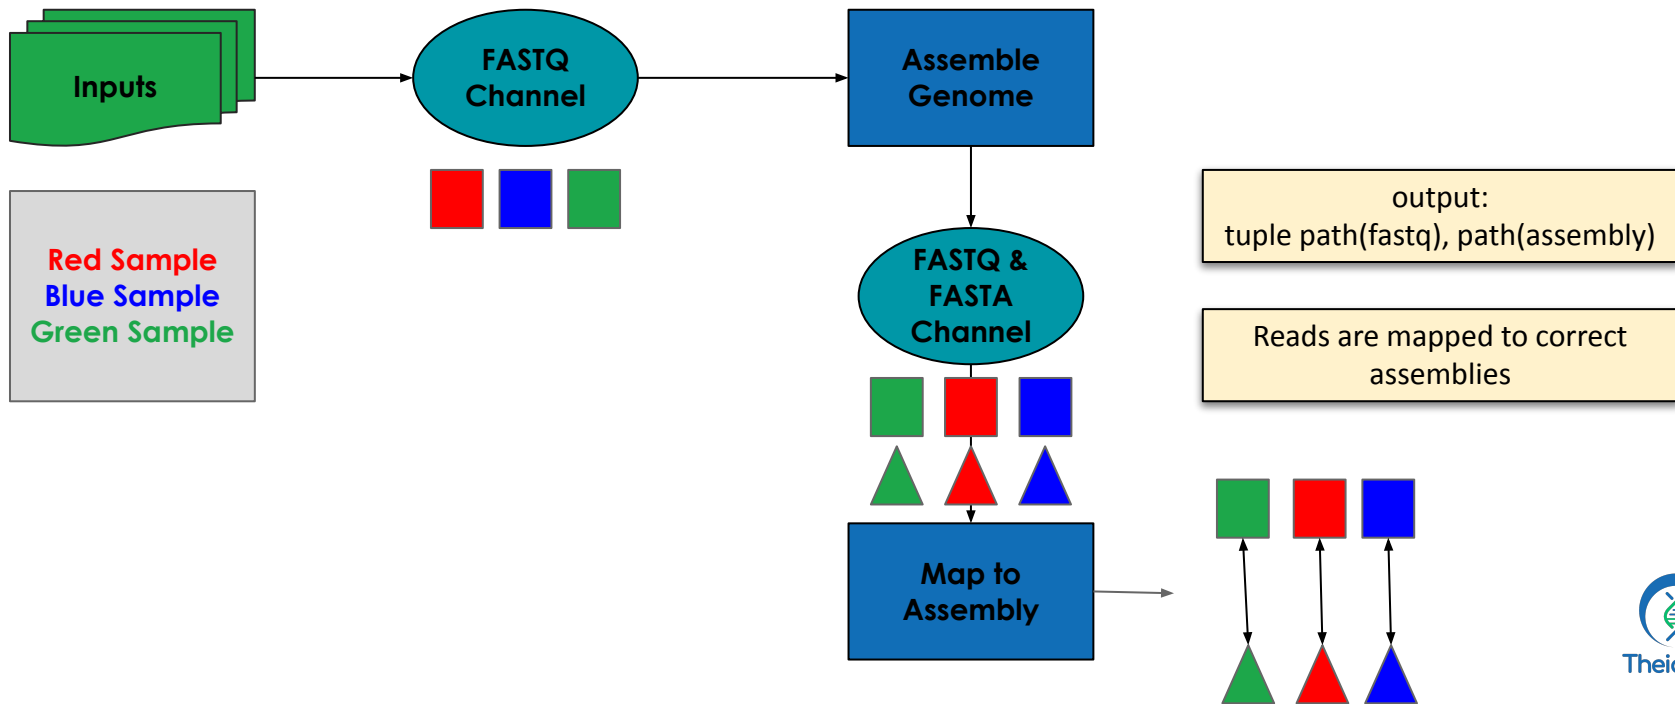
2. Nextflow Files

Nextflow Files

The Nextflow module file
- Modules – A **<u>reusable</u>** Nextflow script with a process definition
- Subworkflows – Multiple modules linked together

> Modules are portable and easily shared between workflows

```
process < name > {

  [ directives ]

  input:
   < process inputs >

  output:
   < process outputs >

  when:
   < condition >


  [script|shell|exec]:
  < user script to be executed >


}
```

Theiagen

Nextflow Files

The Nextflow script file
The `main.nf` file is the central script in a Nextflow workflow.

It defines the structure and execution logic of the pipeline by **orchestrating the processes and connecting them with channels**. This file acts as the "blueprint" for how data flows through the pipeline and how tasks are executed.

```nextflow
#!/usr/bin/env nextflow
nextflow.enable.dsl=2

// Correctly include the process
// definition from the module
include { newModule } from
'./modules/new_module'

workflow {
    // Define the input channel from the
    // user-specified path
    IN_FilePath =
Channel.fromPath(params.input).ifEmpty {
        exit 1, "No file provided with
pattern: ${params.input}"
    }

    // Execute the 'newModule' process
    newModule(IN_FilePath) | view
}
```

Nextflow Files

The Nextflow config file
When a pipeline script is launched, Nextflow looks for configuration files. By default it searches for `nextflow.config` but one can be provided directly via the `-c` `<config-file>` optional parameter.

```
// Define default settings
params {
    input = null
}

// Configure process settings
process {
    cpus = 2
    memory = '2 GB'
    time = '1h'
}

profiles {
    standard {
        // Default profile for local execution
        process.executor = 'local'
    }

    docker {
        process {
            executor = 'docker'
            container = 'ubuntu:jammy'
        }
    }
}
```

Nextflow Files

Putting everything together
This is not obligatory, but following this structure is highly recommended.

The modules live in a `modules` folder where the `main.nf` and `nextflow.config` files are.


```
bin/nextflow/
├── main.nf
├── modules
│   └── fastq_stats.nf
└── nextflow.config
```

3. Running Nextflow Workflows

Running Nextflow Workflows

```
nextflow run main.nf -c nextflow.config --input sample.fastq
```

To run a pipeline:
- All software dependencies must be installed (Nextflow + Docker / Singularity / Conda).
- Configure Nextflow to run on your system (-profile)
- Run the tests for your pipeline in the terminal to confirm everything is working (-profile test)

Core Nextflow command-line options use one (-), whereas pipeline-specific parameters use two (--).

Use -resume to restart pipelines that did not complete. This uses cached results for successful tasks from the previous run, instead of executing all tasks from scratch.

5. Nf-Core
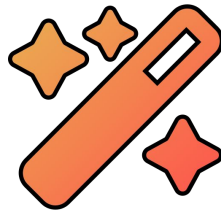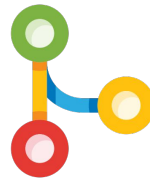
Nf-Core Github Repository

nf-core 🍏

- [nf-core/modules](#) has 500+ ready to use DSL2 modules
  - Rapid prototyping and pipeline building
  - Version controlled, extensive logging, supports Conda, Docker and Singularity

Nextflow DSL2 Module

A community effort to collect a curated set of analysis pipelines built using Nextflow.

6. Awesome Workflows

🔥🐦‍🔥🔥 PHoeNIx: A short-read pipeline for healthcare-associated and antimicrobial resistant pathogens

DOI 10.5281/zenodo.10869898

nextflow DSL2 ≥21.10.3 🐳 run with docker run with singularity

slack StaPH-B #phoenix-dev
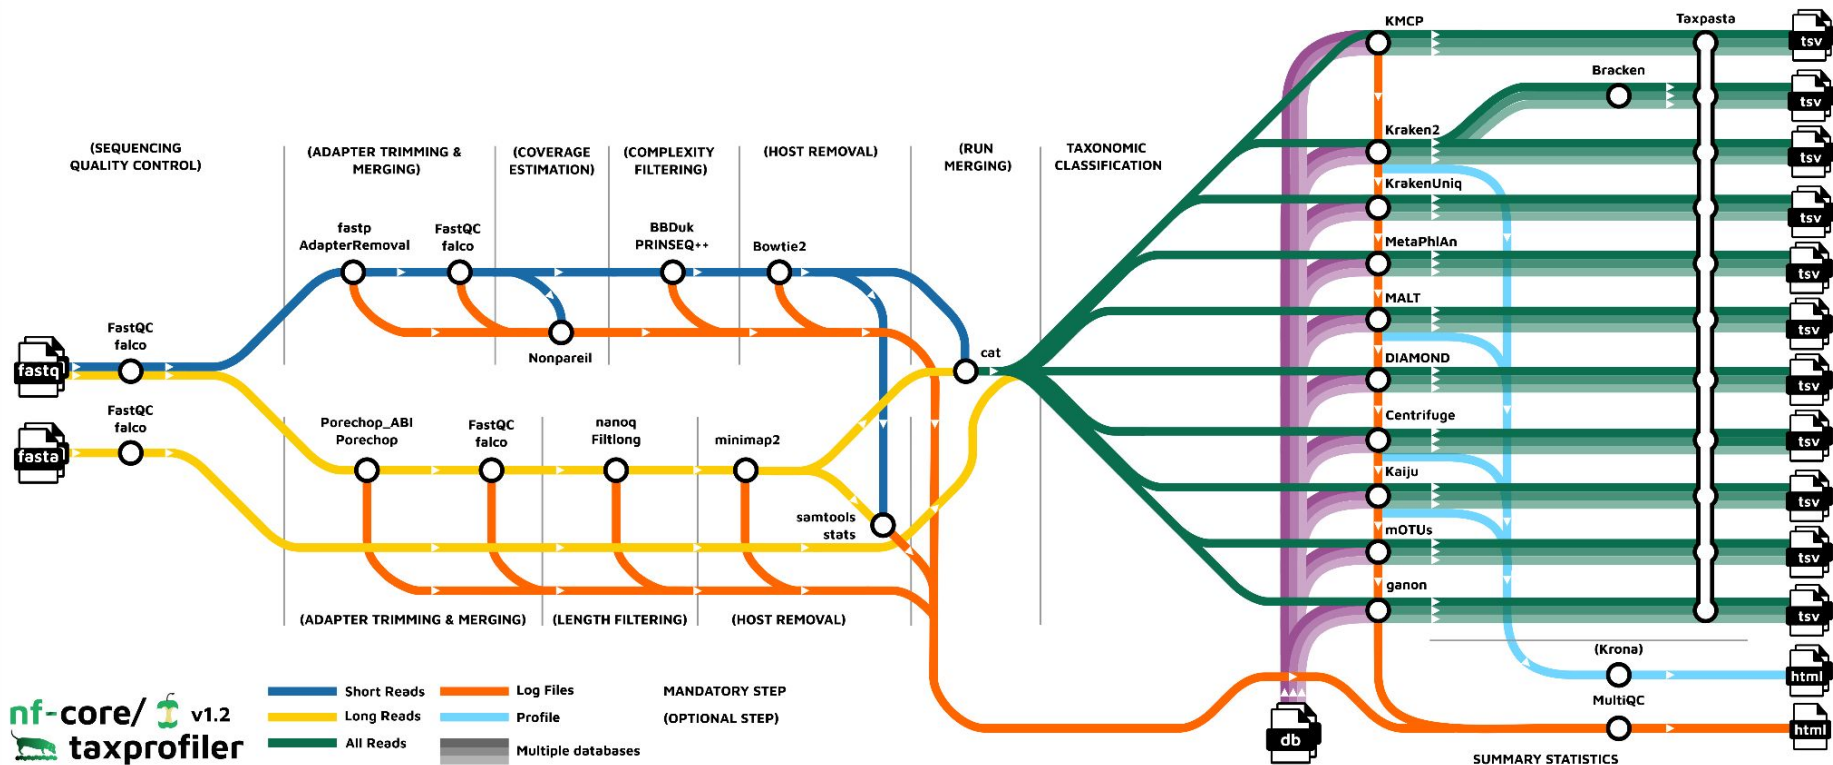
Metrics

Page Hits 24 / 14398

GitHub Clones: 262

For full documentation on the pipeline see the Wiki, but quick start instructions are provided below if you are feeling brave.

https://github.com/CDCgov/phoenix

https://github.com/nf-core/taxprofiler

Hands-On Exercise

Exercise 04: Writing Your First Nextflow Workflow

Exercise Goal

1. **Use GitHub and your dev environment to:**
   a. Install miniconda
   b. Install and test Nextflow
   c. Wrap fastq-peek.sh into a Nextflow workflow